

Experiment 10 : Solving a Markov Decision Process (MDP)

Total Marks: 100

1. Learning Objectives

Upon successful completion of this assignment, students will be able to:

- Define and understand the core components of a Markov Decision Process (MDP):
 - **States (S)**
 - **Actions (A)**
 - **Transition Model (T)**
 - **Reward Function (R)**
 - **Discount Factor (gamma)**
- Translate a "grid world" problem into a formal MDP structure.
- Implement the **Value Iteration** algorithm from scratch using only NumPy.
- Understand and apply the **Bellman Optimality Equation** to calculate state values.
- **Extract an optimal policy** (the best action for each state) from a converged value function.
- Visualize and interpret the resulting value function and policy.
- Analyze how hyperparameters (`gamma` and "living penalty") affect the agent's final behavior.

2. Introduction

This assignment is your first step into Reinforcement Learning. We will focus on the "planning" problem, where we have a perfect *model* of the environment (the MDP) and want to find the *best possible plan* (the optimal policy) before the agent even takes a step.

You will be implementing **Value Iteration**, a classic algorithm that repeatedly applies the Bellman equation to find the true "value" of being in every state.¹ Once we know the value of all states, figuring out the best action is easy: just move to the state with the highest value!

You will solve the "GridWorld" problem, a 3x4 grid with a goal, a "pit" (danger), and walls. Your agent must learn the shortest, safest path to the goal.

3. Prerequisites

Ensure your Python environment has the following libraries installed:

```
pip install numpy matplotlib seaborn
```

4. Experiment Tasks

Task 1: Define the GridWorld (The MDP) (30 Marks)

First, we must define the "rules of the game." You will not use any existing RL libraries (like `gym`). You will define the world yourself.

The world is a 3x4 grid:

- **States (S):** The grid cells. `(0,0)`, `(0,1)`, `(0,2)`, `(0,3)`, etc.
 - **Walls:** There is a wall at `(1,1)`. The agent cannot move into this state.
 - **Terminal States:**
 - **Goal:** `(0,3)` (e.g., a gem)
 - **Pit:** `(1,3)` (e.g., a fire pit)
 - **Actions (A):** The agent can try to move `['up', 'down', 'left', 'right']`.
1. **Define States:** Create a list or set of all valid states (all `(row, col)` tuples, *except* the wall at `(1,1)`).
 2. **Define Rewards (R):** Create a dictionary or function that defines the reward `R(s)` for *being in* a state `s`.
 - **Goal** `(0,3)` : `+1`
 - **Pit** `(1,3)` : `-1`
 - **All other states:** `0.04` (This is a "living penalty" to encourage the agent to find the *shortest* path).
 3. **Define Discount Factor:** Use `gamma = 0.99`.
 4. **Define Transition Model (T):** This is the most important part. You must create a function that defines the *probabilities* of moving. The world is **stochastic** (unpredictable).

- If the agent chooses an action (e.g., `'up'`):
 - **80% chance** it goes in the intended direction (e.g., `'up'`).
 - **10% chance** it slips and goes 90 degrees to the **left** (e.g., `'left'`).
 - **10% chance** it slips and goes 90 degrees to the **right** (e.g., `'right'`).
- **Handling Walls/Boundaries:** If a move (intended or slipped) would land the agent in a wall (like `(1,1)`) or off the grid, the agent **stays in its current state**.²
- **Terminal States:** Once the agent enters a terminal state (Goal or Pit), it stays there and receives no further rewards. (For Value Iteration, we can simplify this: terminal states have a value of 0 and no actions leading out).³

To Implement: Create a helper function `get_next_states(s, a)` that, given a state s and action a , returns a list of `(probability, next_state)` tuples. For example, from `(0,0)`:

`get_next_states((0,0), 'right')` might return:

`[(0.8, (0,1)), (0.1, (0,0)), (0.1, (1,0))]`

(0.8 for 'right', 0.1 for 'up' (slips left, hits wall, stays at $(0,0)$), 0.1 for 'down' (slips right)).

Task 2: Value Iteration Algorithm (From Scratch) (40 Marks)

Now you will implement the algorithm to *solve* the MDP.

1. **Initialize Value Function:** Create your main data structure, V , which will be a dictionary or a 2D NumPy array. $V[s]$ stores the current estimated value of being in state s . Initialize the value of all states to **0.0**.
2. **Implement Value Iteration:**
 - You will loop until the V function converges.
 - Convergence is when the maximum change in V for any state in a single iteration is very small.
 - Use a threshold `theta = 0.0001`.
3. **Run:** Run your `value_iteration` function. It should return the final, converged V table.

Task 3: Policy Extraction (From Scratch) (15 Marks)

The V table tells you how *good* each state is, but not *what to do*. Now, you must extract the optimal policy (P_i) from V .

1. **Create Policy Table:** Create a new table P_i (dictionary or 2D array) to store the best action for each state.

2. One-Step Lookahead:

For each state s :

- Calculate the expected value ($Q(s,a)$) for all four actions, *just like you did in Task 2.*
- Find the action a that gives the **maximum** $Q(s,a)$ value.
- Store this best action (e.g., 'up') in $\text{Pi}[s]$.

3. Return:

Return the final Pi table. This Pi is the optimal policy!

Task 4: Visualization and Analysis (15 Marks)

1. **Visualize Value Function:** Write a simple function to print your V table in a 3x4 grid format. Use `seaborn.heatmap` for a much better visualization.
2. **Visualize Policy:** Write a simple function to print your Pi table in a 3x4 grid, using arrows (\wedge , \vee , $<$, $>$) to represent the actions.

3. Analyze:

- **Question 1:** Run your full pipeline. Print the final V table and the final Pi table. Does the policy make sense? Does it correctly avoid the pit and find the goal?
- **Question 2:** Change the "living penalty" $R(s)$ from 0.04 to 0.0 . Rerun. Does the policy change? Why or why not?
- **Question 3:** Change the "living penalty" $R(s)$ from 0.04 to 0.5 (a high penalty). Rerun. What happens to the policy? Does the agent take a different path? Why?

5. Submission Guidelines

Submit a single `.zip` archive containing:

1. **Source Code:** A single Jupyter Notebook (`.ipynb`) containing all your code, outputs, and visualizations.
2. **PDF Report:** A formal report (`StudentID_Report.pdf`) that includes:
 - **Code Snippets:** Your implementation of the `value_iteration` loop (from Task 2) and your `policy_extraction` function (from Task 3).
 - **Final Results:** The final visualized **Value Table** (as a heatmap) and **Policy Table** (as arrows) for the default parameters ($R=-0.04$).
 - **Analysis:** Your written answers to the three analysis questions in Task 4, explaining the changes in the agent's behavior.